

Norwegian Company Registry API (Brønnøysundregisteret) — VERIFIED

Status: verified against the live API documentation on 2026-08-29

(<https://data.brreg.no/enhetsregisteret/api/docs/index.html>). This replaces the earlier version of this file, which was written from training-time knowledge and flagged as unverified. Corrections and additions from the real docs are marked below.

What it is

Brønnøysundregisteret's **Enhetsregisteret** (Central Coordinating Register for Legal Entities) exposes a free, public, no-auth REST API. Confirmed base URL and shape:

- Base: `https://data.brreg.no/enhetsregisteret/api`
- Single entity: `GET /enheter/{organisasjonsnummer}`
- Search: `GET /enheter?navn={query}` (free-text, ranked; many other filter params available — see full param table in the live docs)
- **Sub-units (subsidiaries/branches):** `GET /underenheter/{organisasjonsnummer}`, each carrying an `overordnetEnhet` field pointing back to its parent org number — ****this is a direct, structured answer to the parent/subsidiary disambiguation problem the challenge brief calls out**, not something that needs to be inferred from crawled text.**
- **Bulk download:** `GET /enheter/lastned` — the entire dataset, gzip-compressed JSON, regenerated nightly around 05:00 (Norway time). For a 411,160-company universe, ****this is very likely a better starting point than 411,160 individual API calls**** — download once, cache locally, and only hit the live per-entity endpoint for refreshes on the daily-random-100 slice.
- ****Incremental updates:** `GET /oppdateringer/enheter?dato=...` or `?oppdateringsid=...` — returns entities changed since a given timestamp or update-sequence-id, with a change type (Ny/Endring/Sletting/Fjernet). Built for exactly the kind of refresh loop `src/storage/snapshot.py` implements — use this instead of re-fetching

every entity to detect what changed.

Confirmed response fields (from the real docs)

Field	Type	Notes
organisasjonsnummer	string	ground-truth ID
navn	string	registered legal name
organisasjonsform.kode / .beskrivelse	object	legal form code + description
hjemmeside	string	confirmed real field — registered website, when present
forretningsadresse	object	business address: adresse[], postnummer, poststed, kommune, kommunenummer, landkode, land
postadresse	object	postal address, same shape — can differ from business address
naeringskode1/2/3	object	industry code(s), kode + beskrivelse
antallAnsatte	number	employee count — directly answers the headcount_band field, no need to infer from LinkedIn
stiftelsesdato	date	founding date
registreringsdatoEnhetsregisteret	date	registration date in the register
konkurs	boolean	bankrupt
underAvvikling	boolean	being wound down
underTvangsavviklingEllerTvangsoppløsning	boolean	forced liquidation/dissolution
slettedato	date	present only if the entity has been deleted from the register
overordnetEnhet (underenheter only)	string	parent org number

Correction to the earlier (unverified) doc

There is **no single status field**. "Active vs. dissolved" must be derived from the combination of `slettedato` (deleted), `konkurs` (bankrupt), and `underAvvikling/underTvangsavviklingEllerTvangsopplosning` (winding down) — richer than the simple active/dissolved binary originally assumed. Update `Entity.status` derivation in `src/resolve/registry.py` to reflect this (currently only checks `slettedato` — good enough as a first pass, but should be extended).

Deleted vs. removed entities

- A **deleted** entity (`slettedato` present) still returns **HTTP 200** with a reduced field set (name, org form, `slettedato`).
- A **removed** entity (e.g. for legal reasons) returns ****HTTP 410 Gone**** with almost no fields. `fetch_registry_record()` should treat 410 the same as any other non-200: return `None`, let the caller degrade to `status="unknown"`.

Rate limits / pagination

- Search results are paginated, default page size 20, and **capped at 10,000 results per query** — $(page+1) * size$ cannot exceed 10,000. Irrelevant for single-entity lookups by org number, but relevant if any fallback logic does name-based search.
- No explicit rate limit documented on individual lookups in what was reviewed — still cache aggressively (see below) rather than assuming unlimited throughput.

Updated recommendation for `src/resolve/`

1. **Download the nightly bulk file once** (`/enheter/lastned`), cache it locally, and resolve the vast majority of the 411,160-company universe from that single download instead of one HTTP call per company — this dramatically reduces both request-budget pressure and wall-clock time versus the per-entity-lookup approach originally sketched.
2. For the **daily random-100 batch specifically**, use live

`/enheter/{orgnr}` lookups (100 calls, trivial against the 2,000-request budget) to ensure freshness, since the bulk file is only regenerated once nightly.

3. Use `/oppdateringer/enheter` to detect what changed since the last bulk download, rather than re-downloading and diffing the entire 411k-row file every refresh cycle.

4. For any resolved parent company, optionally call `/underenheter?overordnetEnheter={orgnr}` to enumerate its subsidiaries — directly useful for the brand/parent/subsidiary disambiguation the brief flags as a common failure mode.

Existing Python client

A community-maintained client, `brreg` (PyPI, Apache 2.0, by Otovo AS), already wraps this API. Worth evaluating as a `requirements.txt` addition instead of hand-rolling the HTTP layer — reduces surface area for bugs in a part of the system everything else depends on. Verify its maintenance status and API coverage before committing to it.

Norwegian legal form codes (`organisasjonsform.kode`) — unchanged, still accurate

See `docs/algorithms/match_algorithm.md` for the AS/ASA/ANS/DA/ENK/KS/BA/SA/NUF suffix table — confirmed consistent with the real `organisasjonsform` codes seen in the live API examples (AS, ENK, BEDR for sub-units, etc.).